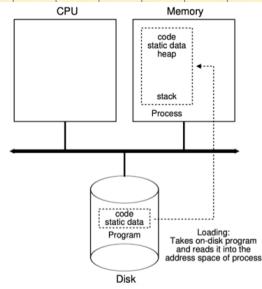
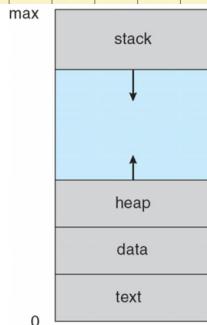


Process

- Process is a program in execution
- A program is a file on the disk
 - Code and static data
- A process is loaded by the OS
 - Code and static data are loaded from the program
 - Heap and stack are created by the OS



- A process is an abstraction of machine states
 - Memory: address space
 - Register:
 - Program Counter (PC) or Instruction Pointer
 - Stack pointer
 - Frame pointer
 - I/O: all files opened by the process



Process Identification

- How can we distinguish processes from one to another?
 - Each process is given a unique ID number, and is called the process ID, or the PID.
 - The system call, getpid(), prints the PID of the calling process.

```
// compile to getpid
#include <stdio.h> // printf()
#include <unistd.h> // getpid()

int main(void) {
    printf("My PID is %d\n", getpid());
}
```

```
$ ./getpid
My PID is 1234
$ ./getpid
My PID is 1235
$ ./getpid
My PID is 1237
```

Process Life Cycle



进程是正在执行的程序

静态：程序（存储在硬盘上）

动态：进程（已创建堆、栈）

进程是机器状态的抽象。

每个进程都有唯一的ID。

PID会不断累加，满时会重新排。

Waiting不能直接返回running

Waiting也称block（等待、阻塞）

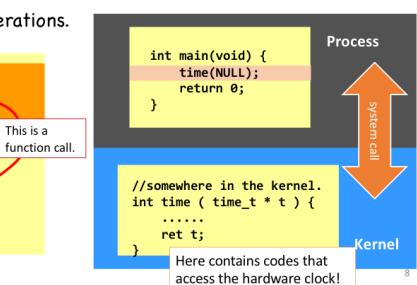
System Call : Process-Kernel Interaction

- System call is a function call.
 - exposed by the kernel.
 - abstraction of kernel operations.

```
int add_function(int a, int b) {
    return (a + b);
}

int main(void) {
    int result;
    result = add_function(a,b);
    return 0;
}

// this is a dummy example...
```

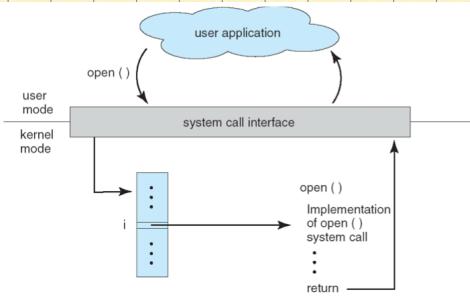


System call 是函数调用

但是 System call 之后不会直接跳到内核态。
而是根据一张表，到内核态调用对应的 System call。

System Call : Call by number

- System call is different from function call
- System call is a call by number



User-mode code from xv6-riscv

```
int main(void) {
    ...
    int fd = open("copyin1", O_CREATE|O_WRONLY);
    ...
    return 0;
}
```

```
/* kernel/syscall.h */
#define SYS_open 15
```

```
/* user/usys.S */
.global open
open:
    li a7, SYS_open
    ecall
    ret
```

Kernel code from xv6-riscv

```
/* kernel/syscall.h */
#define SYS_open 15
```

```
/* kernel/file.c */
uint64 sys_open(void) {
    ...
    return fd;
}
```

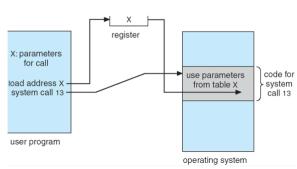
```
/* kernel/syscall.c */
static uint64 (*syscalls[])(void) = {
    ...
    [SYS_open] sys_open,
    ...
};

void syscall(void) {
    struct proc *p = myproc();
    num = p->trapframe->a7;
    p->trapframe->a0 = syscalls[num]();
}
```

System Call : Parameter Passing

- Often, more information is required than the index of desired system call
 - Exact type and amount of information vary according to OS and call
- Three general methods used to pass parameters to the OS
 - **Registers:** pass the parameters in registers
 - In some cases, may be more parameters than registers
 - x86 and risc-v take this approach
 - **Blocks:** Parameters stored in a memory block and address of the block passed as a parameter in a register
 - **Stack:** Parameters placed, or pushed, onto the stack by the program and popped off the stack by the operating system
 - Block and stack methods do not limit the number or length of parameters being passed

Example: parameter passing via blocks

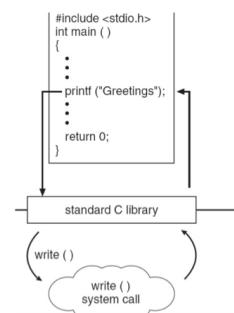


传递参数方式：
寄存器
内存块
栈

System Call vs. Library API Call

- Most operating systems provide standard C library to provide library API calls
 - A layer of indirection for system calls

| Name | System call? |
|--------------------|--------------|
| printf() & scanf() | No |
| malloc() & free() | No |
| fopen() & fclose() | No |
| mkdir() & rmdir() | Yes |
| chown() & chmod() | Yes |

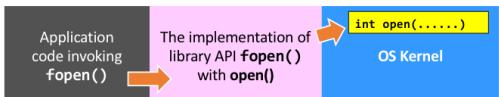


API 不是系统调用

- Take `fopen()` as an example.
 - `fopen()` invokes the system call `open()`.
 - `open()` is too primitive and is not programmer-friendly!

```
Library call   fopen("hello.txt", "w");
```

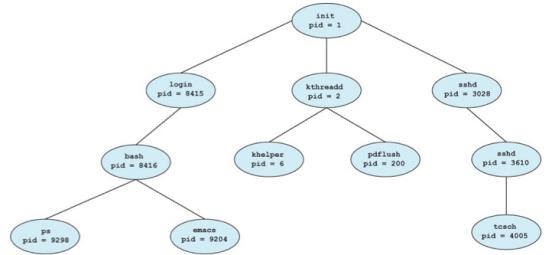
```
System call   open("hello.txt", O_WRONLY | O_CREAT | O_TRUNC, 0666);
```



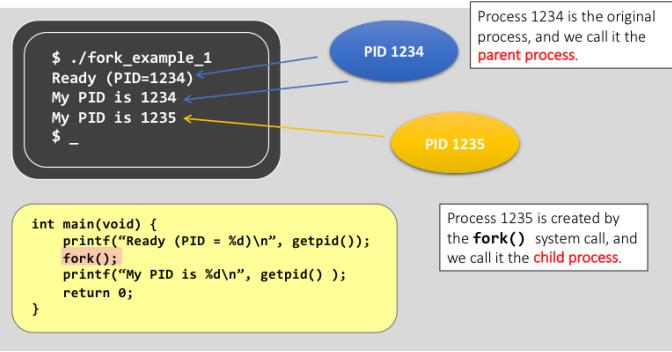
Process Creation

- Parent process creates children processes, which, in turn create other processes, forming a tree of processes
- Generally, process identified and managed via a process identifier (pid)
- Resource sharing
 - Parent and children share all resources
 - Children share subset of parent's resources
 - Parent and child share no resources
- Execution
 - Parent and children execute concurrently
 - Parent waits until children terminate
- Address space
 - Child duplicate of parent
 - Child has a program loaded into it
- UNIX examples
 - fork system call creates new process
 - exec system call used after a fork to replace the process' memory space with a new program

A tree of processes in Linux



Creating Processes with fork() System Call



What do we know so far?

- Both the parent and the child execute **the same program**.
- The child process starts its execution **at the location that fork() is returned, not from the beginning of the program**.

进程创建会形成进程树。

进程由 pid 唯一标识。

资源共享

并发执行

地址空间

父子进程执行相同的程序

子进程从 fork() 返回处开始执行。

```

1 int main(void) {
2     int result;
3     printf("before fork ...\\n");
4     result = fork();
5     printf("result = %d\\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\\n");
9         printf("My PID is %d\\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\\n");
13        printf("My PID is %d\\n", getpid());
14    }
15
16    printf("program terminated.\\n");
17 }

```

\$./fork_example_2
before fork ...

Important

- Both parent and child need to return from fork().
- CPU scheduler decides which to run first.

PID 1234 → fork() → PID 1235

fork()返回后由CPU决定谁先运行。

```

1 int main(void) {
2     int result;
3     printf("before fork ...\\n");
4     result = fork();
5     printf("result = %d\\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\\n");
9         printf("My PID is %d\\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\\n");
13        printf("My PID is %d\\n", getpid());
14    }
15
16    printf("program terminated.\\n");
17 }

```

\$./fork_example_2
before fork ...
result = 1235

Important

- For parent, the return value of fork() is the PID of the created child.

PID 1234 (running) → PID 1235 (waiting)

在父进程中, fork()的返回值是子进程的pid。

```

1 int main(void) {
2     int result;
3     printf("before fork ...\\n");
4     result = fork();
5     printf("result = %d\\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\\n");
9         printf("My PID is %d\\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\\n");
13        printf("My PID is %d\\n", getpid());
14    }
15
16    printf("program terminated.\\n");
17 }

```

\$./fork_example_2
before fork ...
result = 1235
I'm the parent.
My PID is 1234
program terminated.
result = 0

Important

- For child, the return value of fork() is 0.

PID 1234 (stop) → PID 1235 (running)

在子进程中, fork()的返回值是0。

```

1 int main(void) {
2     int result;
3     printf("before fork ...\\n");
4     result = fork();
5     printf("result = %d\\n", result);
6
7     if(result == 0) {
8         printf("I'm the child.\\n");
9         printf("My PID is %d\\n", getpid());
10    }
11    else {
12        printf("I'm the parent.\\n");
13        printf("My PID is %d\\n", getpid());
14    }
15
16    printf("program terminated.\\n");
17 }

```

\$./fork_example_2
before fork ...
result = 1235
I'm the parent.
My PID is 1234
program terminated.
result = 0
I'm the child.
My PID is 1235
program terminated.
\$ _

PID 1234 (stop) → PID 1235 (stop)

fork() System Call

- fork() behaves like "cell division".

• It creates the child process by **cloning** from the parent process, including all user-space data, e.g.,

| Cloned items | Descriptions |
|----------------------------------|-----------------------------------------------------------------------------------------------------|
| Program counter [CPU register] | That's why they both execute from the same line of code after fork() returns. |
| Program code [File & Memory] | They are sharing the same piece of code. |
| Memory | Including local variables, global variables, and dynamically allocated memory. |
| Opened files [Kernel's internal] | If the parent has opened a file "fd", then the child will also have file "fd" opened automatically. |

fork() 会创建新的子进程，父进程的大部分值都会复制到该子进程中，只有少数值不同。

- fork() does not clone the following...

| Distinct items | Parent | Child |
|------------------------|---------------------------|------------------------------------------------|
| Return value of fork() | PID of the child process. | 0 |
| PID | Unchanged. | Different, not necessarily be "Parent PID + 1" |
| Parent process | Unchanged. | Parent. |
| Running time | Cumulated. | Just created, so should be 0. |
| [Advanced] File locks | Unchanged. | None. |

exec() System Call family

- exec() - a member of the exec system call family (exec, execle, execlp, execv, execve, execvp).

```
int main(void) {
    printf("before execl ...\\n");
    execl("/bin/ls", "/bin/ls", NULL);
    printf("after execl ...\\n");
    return 0;
}
```

Arguments of the execl() call

1st argument: the program name, "/bin/ls" in the example.
2nd argument: argument[0] to the program.
3rd argument: argument[1] to the program.

```
int main(void) {
    printf("before execl ...\\n");
    execl("/bin/ls", "/bin/ls", NULL);
    printf("after execl ...\\n");
    return 0;
}
```

What is the output?

The same as the output of running "ls" in the shell.

exec() 的一系列系统调用会启动新的进程，加载不同的程序

- Example #1: run the command "/bin/ls"

| | | |
|------------------------------------|------------------------|-----------------------------------------------------------------------------------------|
| execl("/bin/ls", "/bin/ls", NULL); | | |
| Argument Order | Value in above example | Description |
| 1 | "/bin/ls" | The file that the programmer wants to execute. |
| 2 | "/bin/ls" | When the process switches to "/bin/ls", this string is the program argument[0] . |
| 3 | NULL | This states the end of the program argument list. |

- Example #2: run the command "/bin/ls -l"

| | | |
|------------------------------------------|------------------------|-----------------------------------------------------------------------------------------|
| execl("/bin/ls", "/bin/ls", "-l", NULL); | | |
| Argument Order | Value in above example | Description |
| 1 | "/bin/ls" | The file that the programmer wants to execute. |
| 2 | "/bin/ls" | When the process switches to "/bin/ls", this string is the program argument[0] . |
| 3 | "-l" | When the process switches to "/bin/ls", this string is the program argument[1] . |
| 4 | NULL | This states the end of the program argument list. |

```
int main(void) {
    printf("before execl ...\\n");
    execl("/bin/ls", "/bin/ls", NULL);
    printf("after execl ...\\n");
    return 0;
}
```

The shell prompt appears!

\$. /exec_example
before execl ...
exec_example
exec_example.c
\$ _

The output says:
(1) The gray code block is not reached!
(2) The process is terminated!

WHY IS THAT?

exec() 调用后不会返回原来的程序

```
/* code of program exec_example */
int main(void) {
    printf("before execl ... \n");
    execl("/bin/ls", "/bin/ls", NULL);
    printf("after execl ... \n");
    return 0;
}
```

Address Space
of the process

```
/* Code of program "ls" */
int main(int argc, char ** argv)
{
    .....
    exit(0);
}
```

Address Space
of the process

exec() loads program "ls" into the
memory of this process

```
/* Code of program "ls" */
int main(int argc, char ** argv)
{
    .....
    exit(0);
}
```

Address Space
of the process

The "return" or the "exit()"
statement in "/bin/ls" will terminate
the process...

Therefore, it is certain that the process
cannot go back to the old program!

exec()后进程会进入新程序，不再返回。

exec() Summary

- The process is changing the code that is executing and never returns to the original code.
 - The last two lines of codes are therefore not executed.
- The process that calls an exec* system call will replace user-space info, e.g.,
 - Program Code
 - Memory: local variables, global variables, and dynamically allocated memory;
 - Register value: e.g., the program counter;
- But, the kernel-space info of that process is preserved, including:
 - PID;
 - Process relationship;
 - etc.

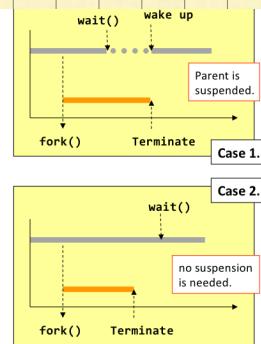
进程中的代码会改变且永不返回原代码。

即会取代用户空间信息

进程的内核空间的信息会保留

Wait()

- wait() suspends the calling process to waiting
- wait() returns when
 - one of its child processes changes from running to terminated.
- Return immediately (i.e., does nothing) if
 - It has no children
 - Or a child terminates before the parent calls wait for



挂起正在运行的进程
当其中一个子进程终止时返回

若①没有子进程 ②在wait()调用之前子进程已停止，则会立刻返回。

Wait() v.s. waitpid()

- wait()
 - Wait for any one of the child processes
 - Detect child termination only
- waitpid()
 - Depending on the parameters, waitpid() will wait for a particular child only
 - Depending on the parameters, waitpid() can detect different status changes of the child (resume/stop by a signal)

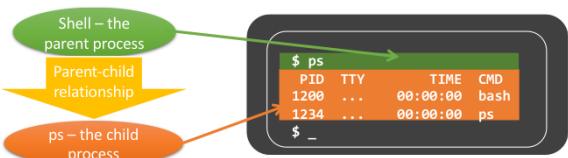
wait()会等待任一子进程，且仅检测子进程停止

waitpid()会等待有特定pid的进程，且检测孩子进程所有的状态变化。

Implement Shell with fork(), exec(), and wait()

- A shell is a CLI

- Bash in linux
- invokes a function fork() to create a new process
- Ask the child process to exec() the target program
- Use wait() to wait until the child process terminates



Processes: Kernel View

Process Control Block (PCB) 进程控制块

Information associated with each process

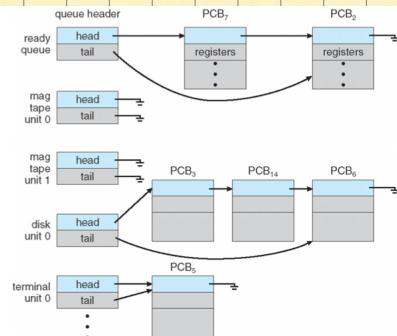
- Process state
- Program counter
- CPU registers
- CPU scheduling information
- Memory-management information
- Accounting information
- I/O status information



Ready Queue And I/O Device Queues

- PCBs are linked in multiple queues

- Ready queue contains all processes in the ready state (to run on this CPU)
- Device queue contains processes waiting for I/O events from this device
- Process may migrate among these queues



PCB 连接在多个队列中。

Ready Queue : 在 ready 状态的进程

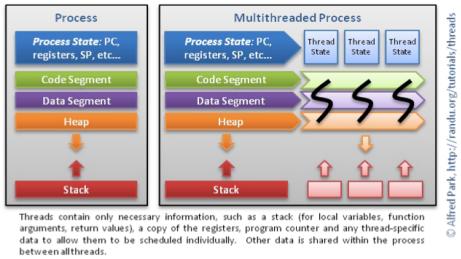
Device Queue : 等待 I/O 的进程

进程可能在多个队列中迁移

Threads 线程

- One process may have more than one threads
 - A single-threaded process performs a single thread of execution
 - A multi-threaded process performs multiple threads of execution "concurrently", thus allowing short response time to user's input even when the main thread is busy
- PCB is extended to include information about each thread

- Single threaded process and multi-threaded process



一个进程中可以有不止一个线程

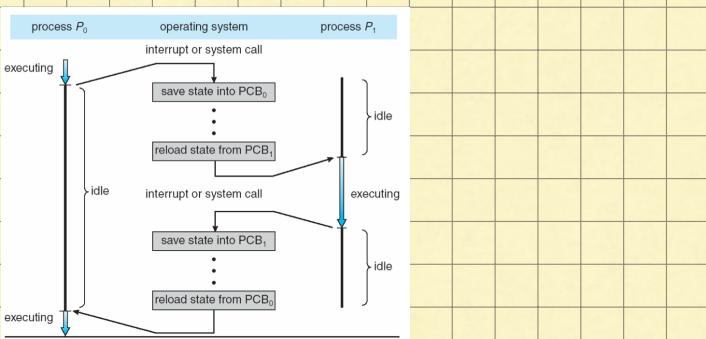
PCB会扩展，包括每个线程的信息

Switching Between Processes

- Once a process runs on a CPU, it only gives back the control of a CPU
 - when it makes a system call
 - when it raises an exception
 - when an interrupt occurs
- What if none of these would happen for a long time?
 - Cooperative scheduling: OS will have to wait
 - Early Macintosh OS, old Alto system
 - Non-cooperative scheduling: timer interrupts
 - Modern operating systems
- When OS kernel regains the control of CPU
 - It first completes the task
 - Serve system call, or
 - Handle interrupt/exception
 - It then decides which process to run next
 - by asking its CPU scheduler
 - How does it make decisions?
 - More about CPU scheduler later
 - It performs a **context switch** if the soon-to-be-executing process is different from the previous one

Context Switch

- During context switch, the system must save the state of the old process and load the saved state for the new process
- Context of a process is represented in the PCB
- The time used to do context switch is an overhead of the system; the system does no useful work while switching
 - Time of context switch depends on hardware support
 - Context switch cannot be too frequent



只有当进入内核态时，CPU才从进程移交给OS
(system call, exception, interrupt)

若长时间不进入内核态：

- ①协作型调度：OS会一直等待
- ②非协作型调度：计时中断

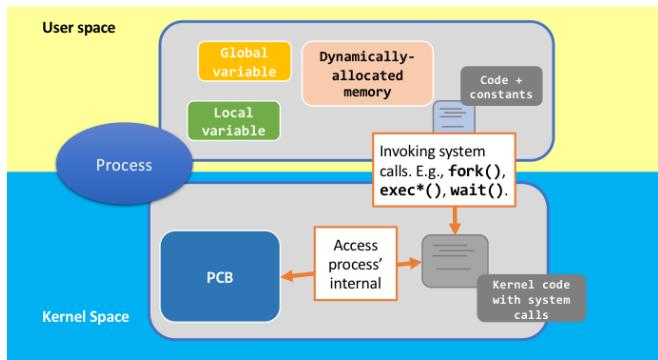
CPU调度器会决定下一个进程

当下一个进程与当前进程不一样时，会发生上下文切换。

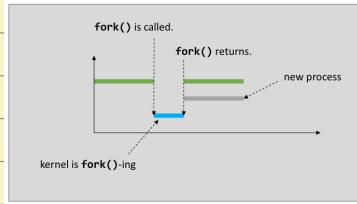
在上下文切换时，系统要保存旧进程的状态，并加载新进程的状态。
进程的上下文信息在PCB里。

上下文切换用时看硬件、不要过于频繁

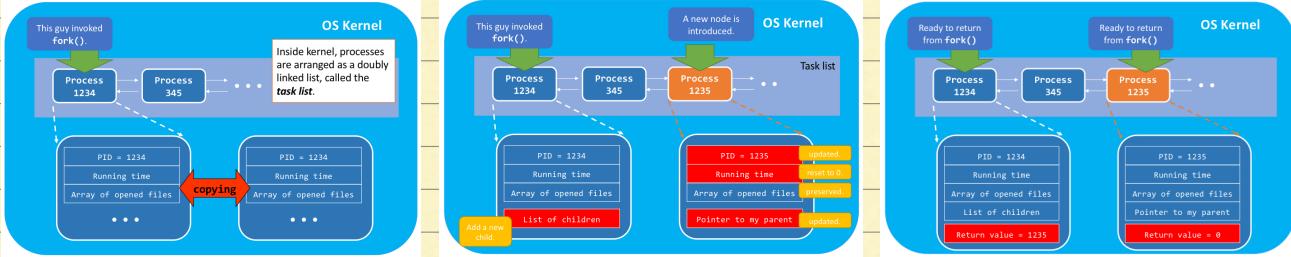
Recall: fork(), exec(), and wait()



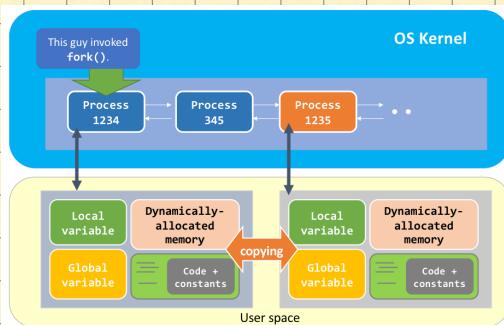
fork() in User Mode



fork(): Kernel View

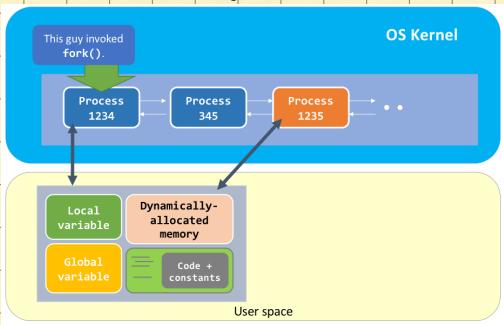


Case 1 : Duplicate Address Space



直接完全复制地址空间

Case 2: Copy on Write



当子进程修改时，才会复制修改的部分

两种实现都有。希望父子进程的地址空间是不一样的。

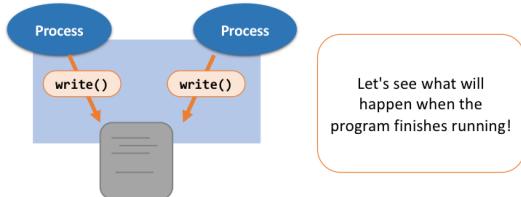
fork(): Opened Files

- Array of opened files contains:

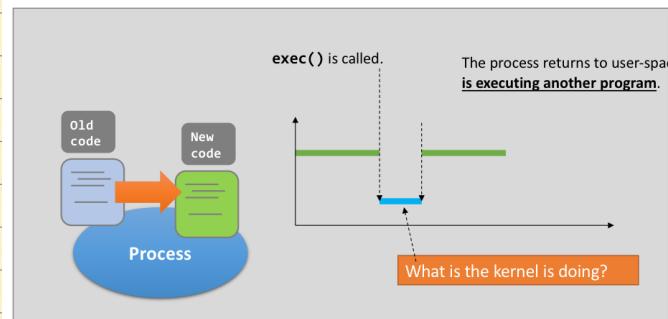
| Array Index | Description |
|-------------|-----------------------------------------------------------|
| 0 | Standard Input Stream; FILE *stdin; |
| 1 | Standard Output Stream; FILE *stdout; |
| 2 | Standard Error Stream; FILE *stderr; |
| 3 or beyond | Storing the files you opened, e.g., fopen(), open(), etc. |

- That's why a parent process shares the same terminal output stream as the child process.

- What if two processes, sharing the same opened file, write to that file together?

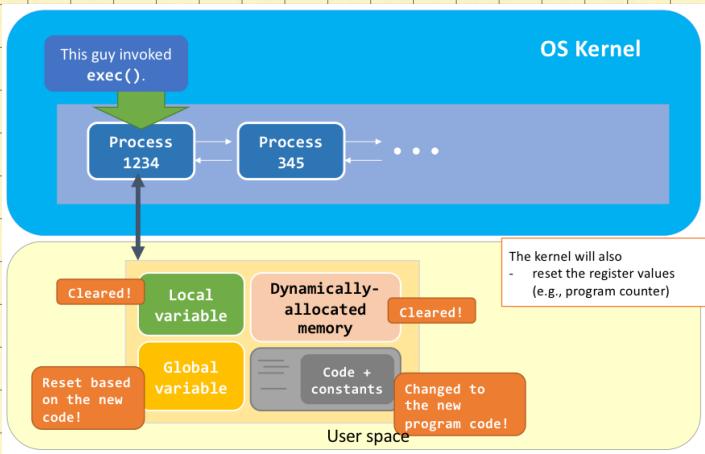


exec() in User Mode



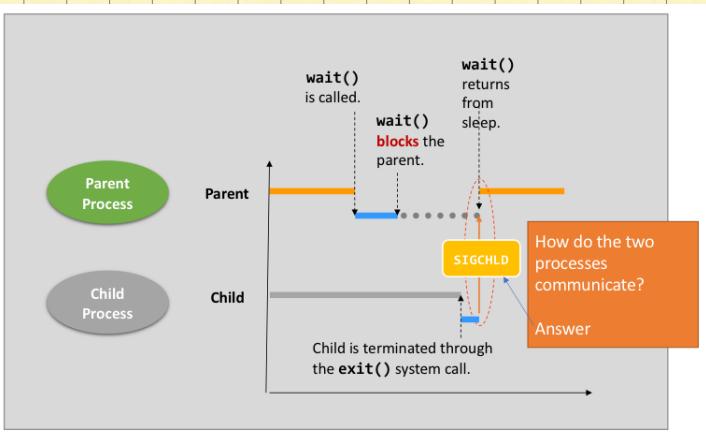
返回 user space 后执行的是另一个程序

exec(): Kernel View

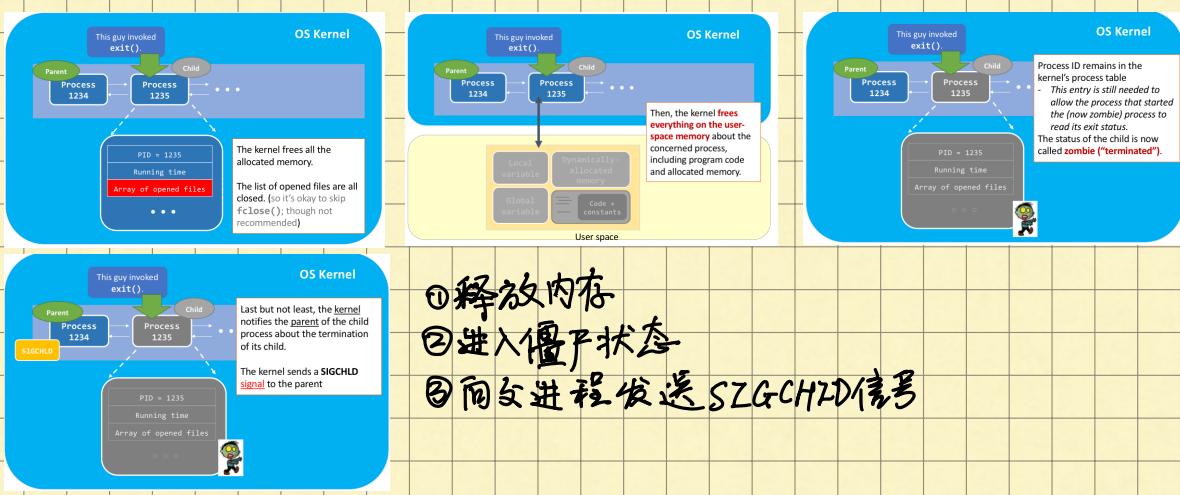


Local variable 与 dynamically allocated 都被清除。
Global 与 code + constant 都变成新程序的

Wait() and exit()

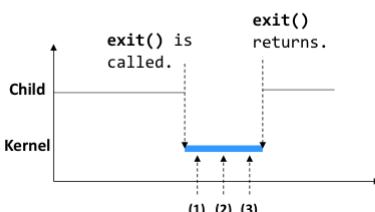


exit() kernel View

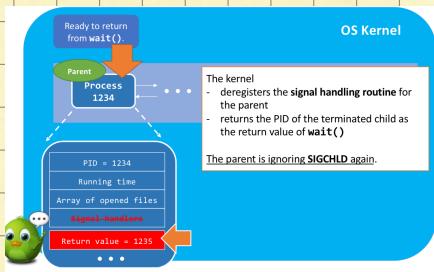
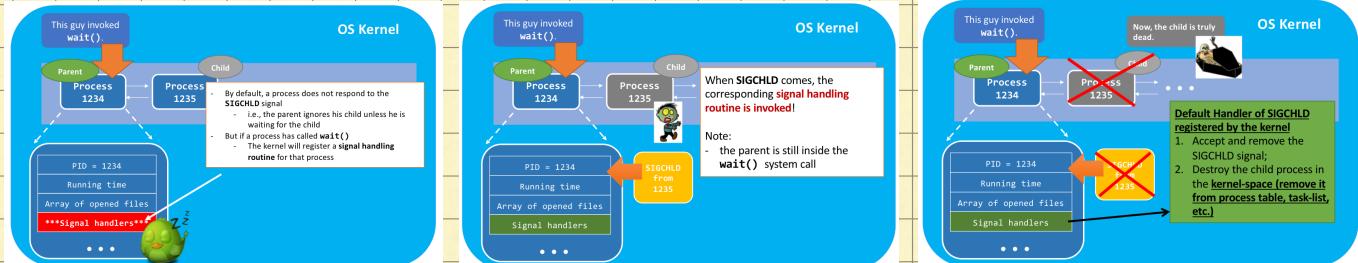


exit() Summary

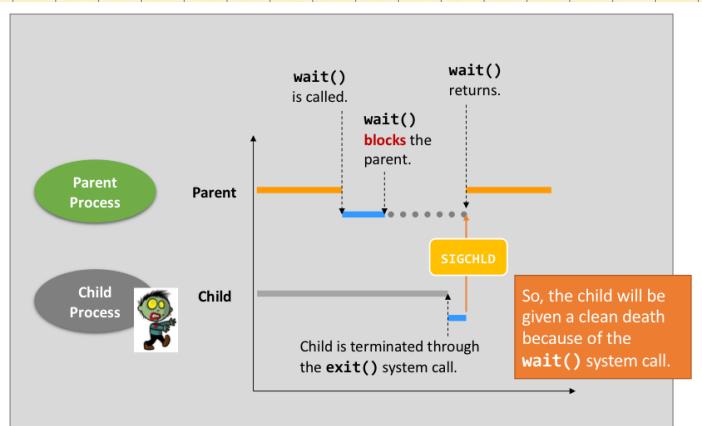
- Step (1) Clean up most of the allocated kernel-space memory (e.g., process's running time info).
- Step (2) Clean up the exit process's user-space memory.
- Step (3) Notify the parent with SIGCHLD.



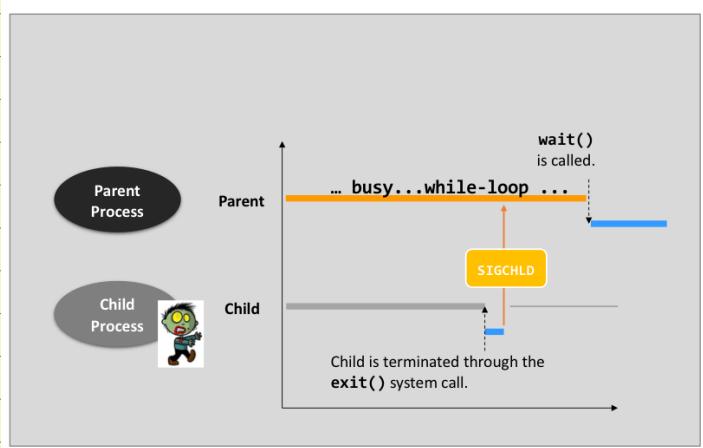
wait() Kernel View



Normal Case



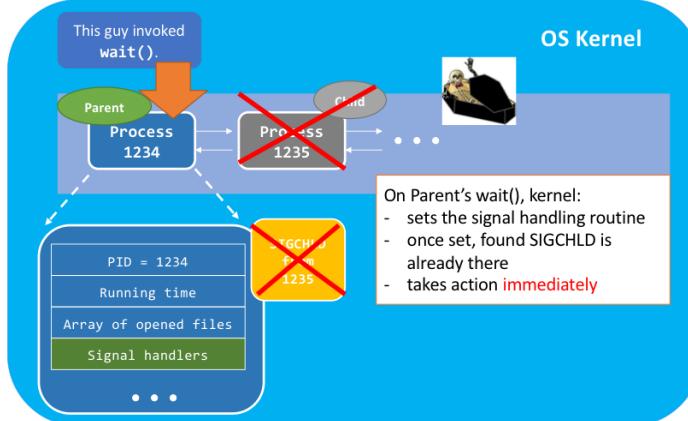
Parent's wait() after Child's exit()



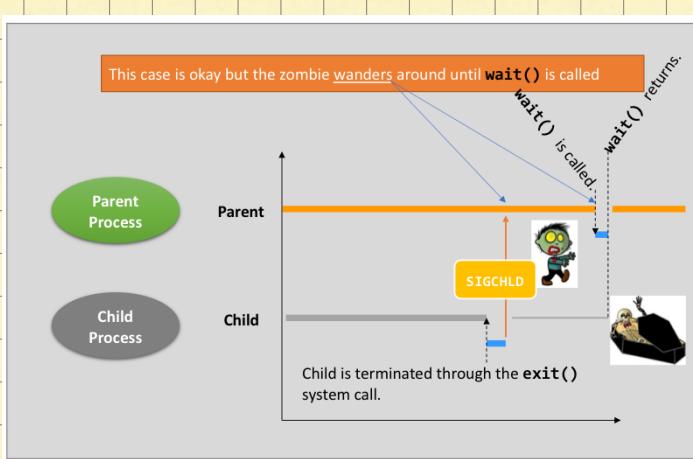
- ① 一般情况下，进程会忽略 SIGCHLD 信号
但若进程调用过 `wait()`，kernel 会为该进程注册 signal handling routine。
- ② 当 SIGCHLD 到来时，对应的 signal handling routine 会被调用
- ③ signal handler 会移除 SIGCHLD 信号，并将 kernel space 的子进程销毁。
- ④ 父进程的 signal handler 会删除，并返回子进程 pid。

当父进程收到 SIGCHLD 时会执行相关操作后结束 `wait()`。

若子进程死亡时父进程仍在执行，SIGCHLD 会在父进程中排队。



当父进程进入wait()时，若发现SIGCHLD已注册有内容，会立刻处理。



Summary of `wait()` & `exit()`

- `exit()` system call turns a process into a zombie when...
 - The process calls `exit()`.
 - The process returns from `main()`.
 - The process terminates abnormally.
 - The kernel knows that the process is terminated abnormally. Hence, the kernel invokes `exit()` for it.
- `wait()` & `waitpid()` reap zombie child processes.
 - It is a must that you should never leave any zombies in the system.
 - `wait()` & `waitpid()` pause the caller until
 - A child terminates/stops, OR
 - The caller receives a signal (i.e., the signal interrupted the `wait()`)
- Linux will label zombie processes as "`<defunct>`".
 - To look for them:

```
$ ps aux | grep defunct
..... 3150 ... [ls] <defunct>
$ -
PID of the process
```

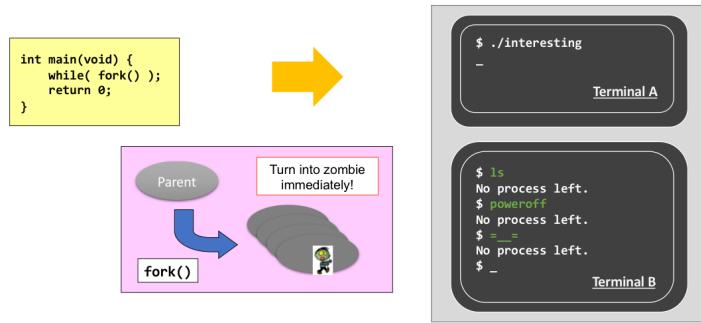
```
1 int main(void)
2 {
3     int pid;
4     if( (pid = fork()) !=0 ) {
5         printf("Look at the status of the child process %d\n", pid);
6         while( getchar() != '\n' ); // "enter" here
7         wait(NULL);
8         printf("Look again!\n");
9         while( getchar() != '\n' ); // "enter" here
10    }
11    return 0;
12 }
```

This program requires you to type "enter" twice before the process terminates.

You are expected to see the status of the child process changes (`ps aux [PID]`) between the 1st and the 2nd "enter".

Using wait() for Resource Management

- It is not only about process execution / suspension...
- It is about system resource management.
 - A zombie takes up a PID;
 - The total number of PIDs are limited;
 - Read the limit: "cat /proc/sys/kernel/pid_max"
 - It is 32,768.
- What will happen if we don't clean up the zombies?



The first process

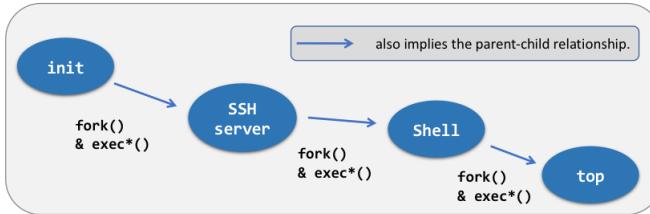
- We now focus on the process-related events.
 - The kernel, while it is booting up, creates the first process - init.
- The "init" process:
 - has PID = 1, and
 - is running the program code "/sbin/init".
- Its first task is to create more processes...
 - Using fork() and exec().

How does uCore
create the first
process?

在内核启动时会创建第一个进程 init。该进程是所有进程的祖先。

A Tree of Processes

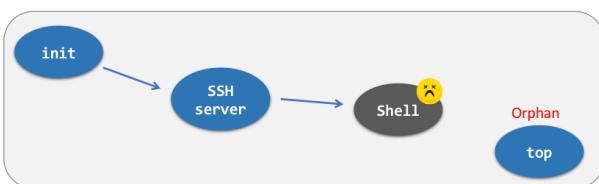
- You can view the tree with the command:



Orphans

- However, termination can happen, at any time and in any place...
 - This is no good because an orphan turns the hierarchy from a tree into a **forest**!
 - Plus, no one would know the termination of the orphan.

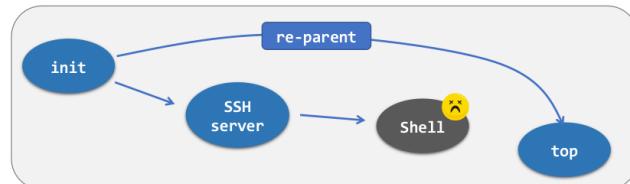
当某进程的父进程先结束了，该进程即为孤儿进程。



Re-parent

- In Linux
 - The "init" process will become the step-mother of all orphans
 - It's called **re-parenting**
- In Windows
 - It maintains a *forest-like process hierarchy*.....

当出现孤儿进程后，该进程的父进程会设成 init。

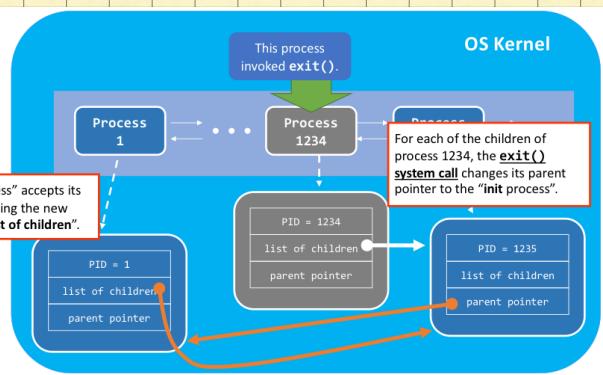
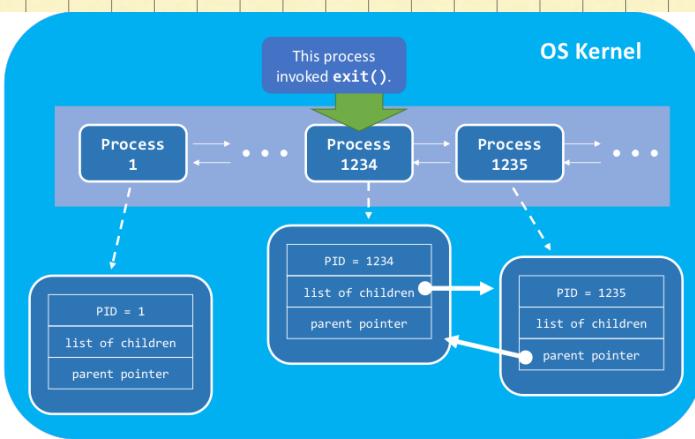
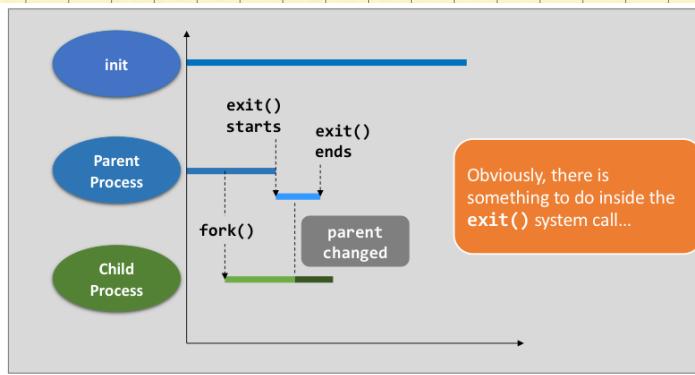


Example

```
1 int main(void) {  
2     int i;  
3     if(fork() == 0) {  
4         for(i = 0; i < 5; i++) {  
5             printf("(%) parent's PID = %d\n",  
6                     getppid(), getpid());  
7             sleep(1);  
8         }  
9     }  
10    else  
11        sleep(1);  
12    printf("(%) bye.\n", getpid());  
13 }
```

getppid() is the system call that returns the parent's PID of the calling process.

```
$ ./reparent  
(1235) parent's PID = 1234  
(1235) parent's PID = 1234  
(1234) bye.  
$ (1235) parent's PID = 1  
(1235) parent's PID = 1  
(1235) parent's PID = 1  
(1235) bye.  
$ -
```



Background Jobs

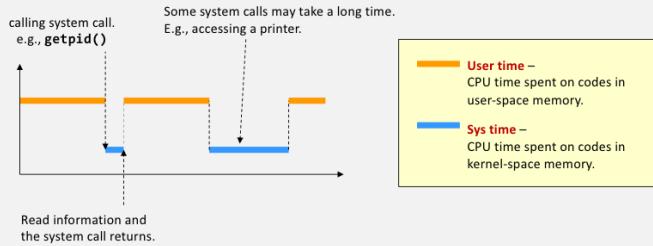
- The re-parenting operation enables something called **background jobs** in Linux
 - It allows a process runs **without a parent terminal/shell**

[Back to home](#)

```
$ ./infinite_loop &
$ exit
[ The shell is gone ]
```

```
$ ps -C infinite_loop
PID TTY
1234 ... ./infinite_loop
$ -
```

Measure Process Time



用户时间：花在用户空间中的CPU时间

系统时间：花在内核空间中的CPU时间

一般来说程序用时是这两个时间相加。

User Time v.s. Sys Time

A terminal window displays the output of the command \$ time ./time_example. The output shows:

```
$ time ./time_example
real 0m0.001s
user 0m0.000s
sys 0m0.000s
$ -
```

Annotations explain the terms:

- Real-time elapsed when "./time_example" terminates.
- The user time of "./time_example".
- The sys time of "./time_example".
- It's possible: real > user + sys
real < user + sys
- Why?
- real>user+sys I/O intensive
real<user+sys multi-core

Below the terminal is a code snippet:

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

A terminal window displays the output of the command \$ time ./time_example. The output shows:

```
$ time ./time_example
real 0m0.001s
user 0m0.000s
sys 0m0.000s
$ -
```

Annotations explain the terms:

- real user sys
- Commented on purpose.

Below the terminal is a code snippet:

```
int main(void) {
    int x = 0;
    for(i = 1; i <= 10000; i++) {
        x = x + i;
        // printf("x = %d\n", x);
    }
    return 0;
}
```

Note: See? Accessing hardware costs the process more time.

I/O密集会经常用户/内核模式切换
多核会统计每个核中的总时间。

The user time and the sys time together define the performance of an application.

- When writing a program, you must consider both the user time and the sys time.
 - E.g., the output of the following two programs are exactly the same. But, their running time is not.

```
#define MAX 100000
int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
#define MAX 1000000
int main(void) {
    int i;
    for(i = 0; i < MAX / 5; i++)
        printf("x\n");
    return 0;
}
```

```
#define MAX 1000000
int main(void) {
    int i;
    for(i = 0; i < MAX; i++)
        printf("x\n");
    return 0;
}
```

```
$ time ./time_example_slow
real 0m1.562s
user 0m0.024s
sys 0m0.108s
$ -
```

```
#define MAX 1000000
int main(void) {
    int i;
    for(i = 0; i < MAX / 5; i++)
        printf("x\n");
    return 0;
}
```

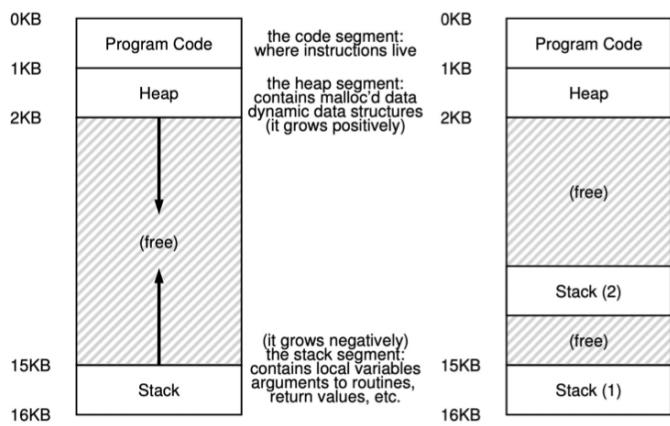
```
$ time ./time_example_fast
real 0m1.293s
user 0m0.012s
sys 0m0.084s
$ -
```

printf()是API，里面有buffer，会在buffer积累到一定程度后才系统调用。

Thread

- Thread is an **abstraction** of the execution of a program
 - A single-threaded program has one point of execution
 - A multi-threaded program has more than one points of execution
- Each thread has its own **private** execution state
 - Program counter and a private set of registers
 - A private stack for thread-local storage
 - CPU switching from one thread to another requires context switch
- Threads in the same process **share** computing resources
 - Address space, files, signals, etc.

Single-Threaded & Multi-Threaded



Thread Pros.

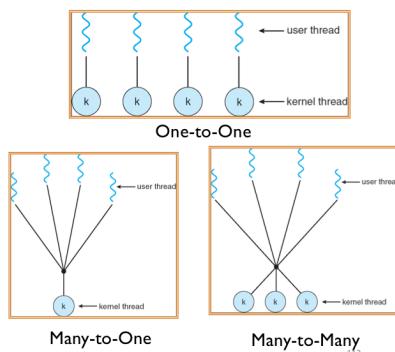
- Increase parallelism
 - One thread per CPU makes better use of multiple CPUs to improve efficiency
- Avoid blocking program progress due to slow I/O
 - Threading enables overlap of I/O with other activities within a single program
 - e.g., many modern server-based applications (web servers, database management systems, and the like) make use of threads
- And allow resource sharing !!!

Thread Implementation

- User-level thread
 - Thread management (e.g., creating, scheduling, termination) done by user-level threads library
 - OS does not know about user-level thread
- Kernel-level thread
 - Thread management done by kernel
 - OS is aware of each kernel-level thread

Thread Models

- One-to-one mapping
 - One user-level thread to one kernel-level thread
- Many-to-one mapping
 - Many user-level thread to one kernel-level thread
- Many-to-many mapping
 - Many user-level thread to many kernel-level thread



线程是程序执行的抽象。

线程有私有执行空间

同一进程中的线程共享计算资源

提高并行

避免因I/O的阻塞

资源共享

用户级线程：

OS并不知道user-level thread

内核级线程：

OS会关心每一个kernel-level thread

Map: 将用户级线程映射到内核级线程。

Pros & Cons

- Many-to-one mapping

- Pros: context switch between threads is cheap
- Cons: When one thread blocks on I/O, all threads block

- One-to-one mapping

- Pros: Every thread can run or block independently
- Cons: Need to make a crossing into kernel mode to schedule

- Many-to-many mapping

- Many user-level threads multiplexed on less or equal number of kernel-level threads
- Pros: best of the two worlds, more flexible
- Cons: difficult to implement